

---

Due date: 25 March 2024, 23:59

---

## 1. Computing $\pi$ with OpenMP [20 points]

1. *Parallelize the serial implementation using OpenMP. Implement two different versions using both the critical directive and the reduction clause.*

In this exercise, we will use the critical directive and the reduction clause to parallelize the serial implementation of the computation of  $\pi$ . The **critical** directive is a synchronization Construct that specifies a region of code that must be executed by only one thread at a time. [2] You have to use the following structure for the parallel region:

```
#pragma omp for parallel shared(x)
for (...) {
    #pragma omp critical
    ...
}
```

On the other hand, the **reduction** is a Data Scope Attribute Clause. It performs a reduction on the variables that appear in its list. A private copy for each list variable is created for each thread. At the end of the reduction, the reduction variable is applied to all private copies of the shared variable, and the final result is written to the global shared variable. [1] You implement the following structure for the region you want to parallelize:

```
#pragma omp parallel for default(shared) private(i, x) reduction(+ : sum)
```

2. *Perform a weak and strong scaling study of your implementations and interpret the results in your report. In particular, discuss the observed differences between the version using critical directive or the reduction clause.*

First we assume a fixed problem, which is to be solved by  $N$  workers. This is called **strong scaling** because the amount of work stays constant no matter how many workers are used. Here the goal of parallelization is minimization of time to solution for a given problem, where the total amount of work is  $s + p$ . This is also known as Amdahl's law. To plot the strong scaling efficiency, we use the following formula for the speedup  $S$ ,

$$S = \frac{T(1)}{T(N)}, \quad (1)$$

where  $T(1)$  is the runtime for the serial code, running in one processor, while  $T(N)$  is the runtime for the parallel code with  $N$  workers. For our implementation of the code to calculate  $\pi$ , we present the results of the strong scaling efficiency in Figure 1. From the figure, we can see that the efficiency of the parallel code with the reduction clause is quite similar than

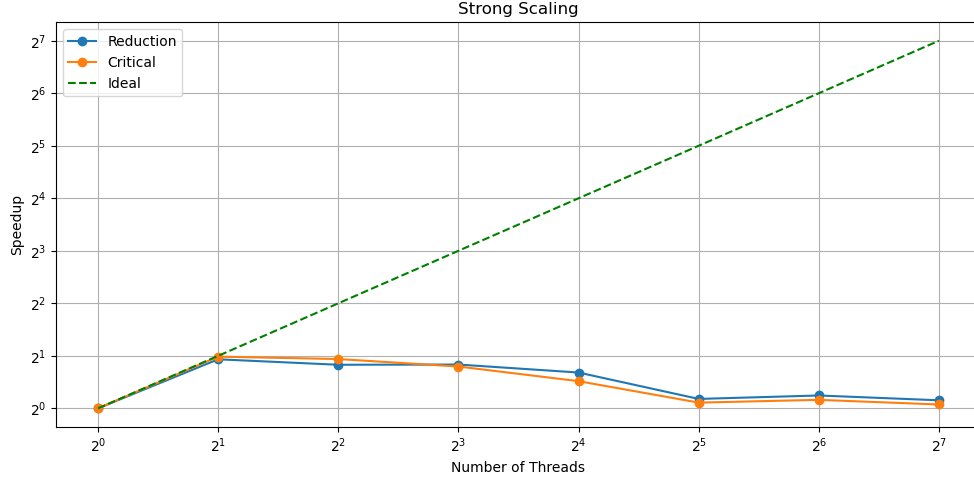


Figure 1: Strong scaling efficiency

that with the critical directive. The analysis was done taking the mean of 100 trials for each number of threads. We can also see that the results are far from the ideal efficiency of slope one, but this is expected and often the case for parallelization efforts. The discrepancy can be attributed to the overhead of communication and synchronization between threads, which is not present in the serial code. We could also try making the problem size larger to see if the speedup improves.

If time to solution is not the primary objective because larger problem sizes (with limited memory) are of interest, it is appropriate to perform what is known as **weak scaling**. In this analysis, we scale the problem size with some power of  $N$  and analyze the difference in efficiency. We use the following formula to calculate efficiency  $E$ ,

$$E = \frac{T(1)}{p * T(N)}. \quad (2)$$

The difference between strong and weak scaling is that in the former the problem size is fixed, while in the latter the problem size is scaled with the number of workers. For our problem, we present the results of the weak scaling efficiency in Figure 2. Similarly to the strong scaling,




Figure 2: Weak scaling efficiency

we can see that both implementations produce similar results. The same arguments can be made for the discrepancy between the ideal efficiency and the actual efficiency.

## 2. The Mandelbrot set using OpenMP [20 points]

This exercise is about parallelizing the computation of the Mandelbrot set using the `OpenMP` directive. The Mandelbrot set is a set of complex numbers  $c$  for which the function  $f_c(z) = z^2 + c$  does not diverge when iterated from  $z = 0$ . Bellow is the code snippet that shows the sequential implementation on the provided template.

[illegible]

The results of the sequential implementation are shown in Figure 3. For all calculations, we used the dimensions  $4096 \times 4096 = 16777216$  pixels. To parallelize code, we can use the **OpenMP** directive `parallel for`. This directive will distribute the iterations of the outer loop among the threads.

```
#pragma omp parallel shared(pPng, fDeltaX, fDeltaY, nTotalIterationsCount,  
    num_threads) private(j, i, x, y, y2, x2, cx, cy, n, c)
```

We also add the directive `#pragma omp for` for the outer loop. Other changes include changing the loop index variables to private, and the `cy` and `cx`, that are the real and imaginary parts of the complex number  $c$ , to be calculated from the loop index variables.

```
cy = fDeltaY * j + MIN_Y;  
cx = fDeltaX * i + MIN_X;
```




Figure 3: Mandelbrot set sequential implementation

The resulting image is shown in Figure 4. We also performed a strong scaling study of the parallel implementation. The results are shown in Figure 5.

Threads	Total time	Avg. time/pixel	Avg. time/iter	Iter/sec	MFlop/s	N iter
Serial	323.388	1.9275E-05	2.85E-09	3.51E+08	2810.86	113624527400
1	334.444	1.9934E-05	2.94E-09	3.40E+08	2718.6	113652339001
2	167.27	9.9700E-06	1.47E-09	6.79E+08	5434.83	113635259482
4	161.266	9.6122E-06	1.42E-09	7.05E+08	5637.06	113633417598
8	109.512	6.5274E-06	9.64E-10	1.04E+09	8300.58	113626596677
16	68.0375	4.0553E-06	5.99E-10	1.67E+09	13355.3	113582813215
32	52.095	3.1051E-06	4.59E-10	2.18E+09	17442.4	113582841475
64	46.8211	2.7907E-06	4.12E-10	2.43E+09	19411	113605732438
128	45.9884	2.7411E-06	4.05E-10	2.47E+09	19762.9	113607727564

Table 1: Performance Metrics. Time is presented in seconds. “Iter” is short for iterations.

We can see some discrepancies in the number of iterations for the different number of threads. This is some kind of numerical instability or incorrect implementation of race conditions. We also see that there are more points than there should be on the x-axis that do not appear on the sequential implementation. We had a lengthy discussion on this issue, but we were not able to find a solution. It may be related to the truncation of the floating point numbers.

On the other hand, we can see that the performance metrics are quite good. The number of iterations is quite large and this allows for the parallelization to be effective. We can see that the number of MFlop/s is increasing with the number of threads, which is expected and a good sign of strong scaling.

### 3. Bug hunt [10 points]

1. **Bug 1** This code has a compilation error and the output is:

```
omp_bug1.c:25:3: error: statement after '#pragma omp parallel for' must be a
    for loop
    {
```




Figure 4: Mandelbrot set parallel implementation with 128 threads.




Figure 5: Mandelbrot set strong scaling efficiency

1 error generated.

The error involves an incorrect placing of the `parallel for` directive. It must be placed before the for loop but in the code it has an extra line getting the number of threads before the for loop section starts. The correct implementation involves moving this line inside the for loop and removing the extra curly brackets.

2. **Bug 2** The second code has a bug at runtime. When I change the `OMP_NUM_THREADS=4` and then run the code, I get this output:

```
Thread 2 is starting...
Number of threads = 4
Thread 0 is starting...
Thread 1 is starting...
Thread 3 is starting...
Thread 3 is done! Total= 0.000000e+00
```

The problem is that the `total` variable is not being updated correctly. The solution is to declare the total variable as a shared variable by adding `shared(total)` to the `parallel` directive. Also the location of setting the total to zero must be before the parallel structure starts. SO that it is not being reinitialized iwth each thread. Finally, `reduction(+ : total)` or a similar class is needed to update the total variable correctly. Additionally, the print statement is only useful if you want to know that the program has ended but not for each thread.

3. **Bug 3** Code three also has a bug at runtime. The idea behind the code is to perform two sections on different threads. When running the program with 2 threads, it performs the action and ends the program correctky. However, when running it with 4 threads, I get the whole output but then the program hangs. What I did was to remove the `#pragma omp barrier` directive in the function `print_results`. This directive is not needed because the `#pragma omp sections` directive already has a barrier at the end of the section. The barrier directive is causing the program to hang because it is waiting for all threads to reach the barrier, but it is never reached because the program is already done.
4. **Bug 4** Code number four causes a segmentation fault. Apparently, the error is related to a memory allocation problem. The solution can be done two ways, by setting the `export OMP_STACKSIZE="9M"` manually or allocating memory dynamically using `malloc` and `free`. The second option is the best solution because it is more portable and does not depend on the system settings, but the first option is also valid and does not require any code changes. We choose 9M because we know it is an array of 1048 elements of type double, so we need at least  $1048 * 8 = 8384$  bytes. We add a little more to be safe.
5. **Bug 5** Code number five creates another runtime error that keeps the program hanging. The issue appears to be related to the locks set between the threads. The error occurs because the locks are not being set and released properly. The solution is to set and unset corrsponding lockes referring to the initialization of both arrays and then the same but calling for the other arrays lock for the addition of both of them. This will prevent the deadlock by ensuring that both threads are not waiting indefinitely for each other to release locks. The implementation of only one section is exemplified below, and the structure is the same for the other section, but with the locks inverted.

```
#pragma omp section
{
    printf("Thread %d initializing a[]\n", tid);
    omp_set_lock(&locka);
    for (i = 0; i < N; i++)
        a[i] = i * DELTA;
    printf("Thread %d done with a[]\n", tid);
    omp_unset_lock(&locka);

    omp_set_lock(&lockb);
    printf("Thread %d adding a[] to b[]\n", tid);
    for (i = 0; i < N; i++)
        b[i] += a[i];
    omp_unset_lock(&lockb);
    printf("Thread %d done adding a[] to b[]\n", tid);
}
```

## 4. Parallel histogram calculation using OpenMP [15 points]

In this exercise, we will parallelize the computation of a histogram using the `OpenMP` directive. The parallelization inlcudes an adaptation of the logic behind a serial implementation to have a private

copy of the histogram for each thread. The private copies are then combined into a shared histogram at the end of the parallel region. The code snippet below shows the parallel implementation of the histogram calculation.

```
#pragma omp parallel shared(num_threads, dist)
{
    num_threads = omp_get_num_threads();
    // Private copy for each thread
    long private_dist[BINS] = {0};

#pragma omp for
    for (long i = 0; i < VEC_SIZE; ++i)
    {
        private_dist[vec[i]]++;
    }

    // Combine private_dist arrays into dist
    for (int i = 0; i < BINS; ++i)
    {
#pragma omp atomic
        dist[i] += private_dist[i];
    }
}
```

Reported are the runtimes of the original (serial) code, as well as the 1-thread and N-thread parallel implementations. The results are shown in Table 2.

Threads	Runtime (s)
Serial	1.50307
1	2.81964
2	1.94883
4	0.882659
8	0.543278
16	0.65781
32	0.279733
64	0.19542
128	0.21948

Table 2: Runtimes of the original (serial) code, as well as the 1-thread and N-thread parallel implementations.

The results of the strong scaling of the parallel implementation are shown in Figure 6. From the results, we can see that the parallel implementation is not as efficient for the 1-thread case as the serial implementation. This is expected because of the change in logic using creating a private data structure for each independent thread. However, the parallel implementation becomes more efficient for the 4-thread case and beyond. The serial case is by itself a good implementation due to the small problem size, and the use of optimization flags in the compilation of the code. We can argue that as the problem size increases, the parallel implementation will be more efficient than the serial implementation. Additionally, only one run was performed for each number of threads, so the results may not be very accurate. More runs would be beneficial to get a better idea of the efficiency of the serial and parallel code.

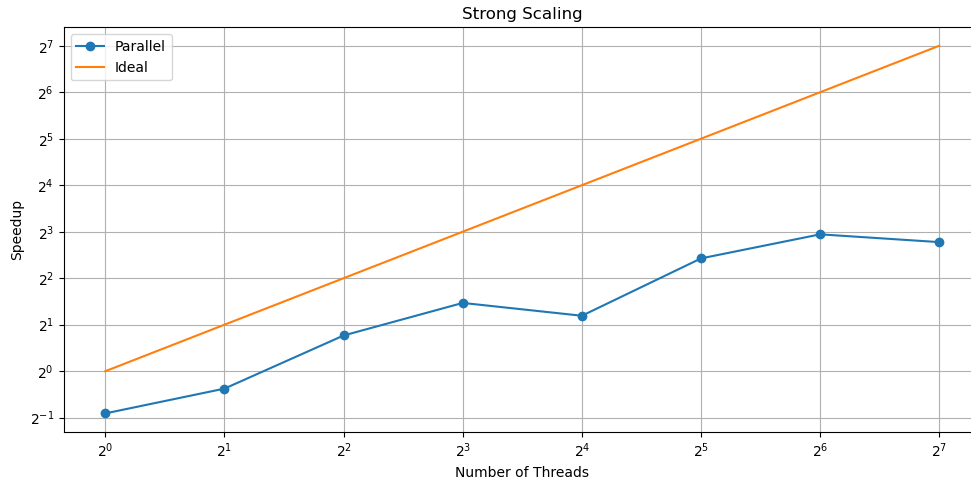


Figure 6: Histogram calculation strong scaling efficiency

## 5. Parallel loop dependencies with OpenMP [15 points]

This exercise is about parallelizing a loop with dependencies using the `firstprivate` and `lastprivate`. The trick behind this exercise is to introduce the power function to calculate the  $S_n$  value at the beginning of each iteration. The code snippet below shows the parallel implementation of the loop with dependencies.

```
#pragma omp parallel shared(opt, N, up) private(n, i)
{
    i = 0;
    #pragma omp for firstprivate(Sn) lastprivate(Sn)
    for (n = 0; n <= N; ++n)
    {
        if (i == 0)
        {
            Sn = Sn * pow(up, n);
            i += 1;
        }
        opt[n] = Sn;
        Sn *= up;
    }
}
```

The Table 3 shows the results of the parallel implementation in terms of the runtimes as well as the final result of the  $S_n$  value and the squared norm of the *opt* vector. The results are quite similar but we can see some discrepancies in the results. These can be attributed to the floating point arithmetic and the parallelization of the code. The results of the strong scaling of the parallel implementation are shown in Figure 7. We can see that the parallel implementation is quite efficient in reducing the total runtime of the code and follows an expected trend common in parallel computing.

## 6. Quicksort using OpenMP tasks [20 points]

The final task of the report involves parallelizing the quicksort algorithm using the OpenMP tasks directive. For my implementation, I used the `task` directive to create tasks for the recursive calls of the quicksort. I also used the `single` directive to create a single task per thread to handle the partitioning of the array. The code snippets below shows the parallel implementation of the quicksort algorithm.



Threads	Parallel Runtime	Final Result $S_n$	Result $\ opt\ _2^2$
Serial	8.232878	485165097.62511122	5884629305179574
1	8.144716	485165097.62511122	5884629305179574
2	4.018959	485165097.62535083	5884629305185405
4	2.301102	485165097.62508273	5884629305178888
8	1.933245	485165097.62554055	5884629305189930
16	1.392649	485165097.62517869	5884629305181860
32	1.173677	485165097.62515235	5884629305181374
64	1.332998	485165097.62513387	5884629305180411
128	1.109134	485165097.62501615	5884629305179745

Table 3: Parallel Runtimes and Results. The time is shown in seconds.




Figure 7: Loop dependencies strong scaling efficiency

```
#pragma omp task
    quicksort(data, right);
#pragma omp task
    quicksort(&(data[left]), length - left);
```

```
#pragma omp parallel
#pragma omp single
quicksort(data, length);
```

Figure 8 shows the results of the strong scaling of the parallel implementation of the quicksort algorithm for different array sizes. We can see some interesting results for the different sizes, but the general trend is that the parallel implementation works well on the lower number of threads than on the higher number of threads. The implementation is quite efficient for the smaller array sizes, but it seems to be even better for the larger array sizes. This is expected because the larger array sizes require more work to be done per thread. We can also see an improvement beyond the ideal line on the largest array size,  $1e + 09$ , but we can argue that since this was only analyzed over one run, that the time may not be very accurate. For the smaller array size, we can see that the parallel implementation becomes less efficient as the number of threads increases, even less efficient than the run on a single thread in some cases. A good consideration to take into account is that the code is the recursive nature of the algorithm; the work that is done in the last levels of the recursion is very small, and therefore using so many threads is not beneficial. The overhead of synchronization and communication between threads is not worth the small amount of work that is done. This is why we see the results for the larger array sizes not being as good as for the smaller




Figure 8: Quicksort parallel implementation

array sizes. In conclusion, the results are quite good and show that the parallel implementation is quite efficient in reducing the total runtime of the code for a limited number of threads, reaching results close to ideal up to 8 threads.

## References

- [1] Lawrence Livermore National Laboratory. OpenMP Directives: Data Scope Attribute Clauses: REDUCTION Clause. [https://hpc-tutorials.llnl.gov/openmp/reduction\\_clause/](https://hpc-tutorials.llnl.gov/openmp/reduction_clause/), 2024. [Online; accessed March 15, 2024].
- [2] Lawrence Livermore National Laboratory. OpenMP Directives: Synchronization Constructs: CRITICAL Directive. [https://hpc-tutorials.llnl.gov/openmp/critical\\_directive/](https://hpc-tutorials.llnl.gov/openmp/critical_directive/), 2024. [Online; accessed March 15, 2024].